


▶ MACHINE LEARNING ◀

Joshua Savoie

AGENDA

- An Introduction
- First Steps
- MNIST Database
- Population-Based Training
- Demo
- Recap
- Going Forward

Neural Networks



From the
ground up



FIRST STEPS



RESEARCH AND PLANNING



- 2 weeks of heavy research + light curiosity in the prior 2 months.
- Sourced many scholarly articles, blogs, and videos.
- Created a general structural layout for the neural network.
- This all allowed the foundations for the neural network to be coded within a few days.

POISONOUS FRUIT

- The task: determine whether a fruit was poisonous or not.
- The fruit has three characteristics that act as inputs to the network:
 - one (1): Length of the spikes;
 - two (2): Density of the spots;
 - three (3): The size of the fruit itself.
- This was solved in the matter of a minute at max by the AI.

CODE STRUCTURE


892 Lines of Code



▶ MNIST DATABASE

"HELLO WORLD"

- The MNIST database is considered the "hello world" of machine learning.
- The task: determine the number (0 - 9) represented by a handwritten digit.
- 784 inputs (28x28 pixels) with each input ranging on a scale of 0.0 - 1.0, 0 being black, 1 being white.




EXAMPLES




▶ POPULATION-BASED TRAINING

A training method developed by Google DeepMind.

POPULATION BASED TRAINING

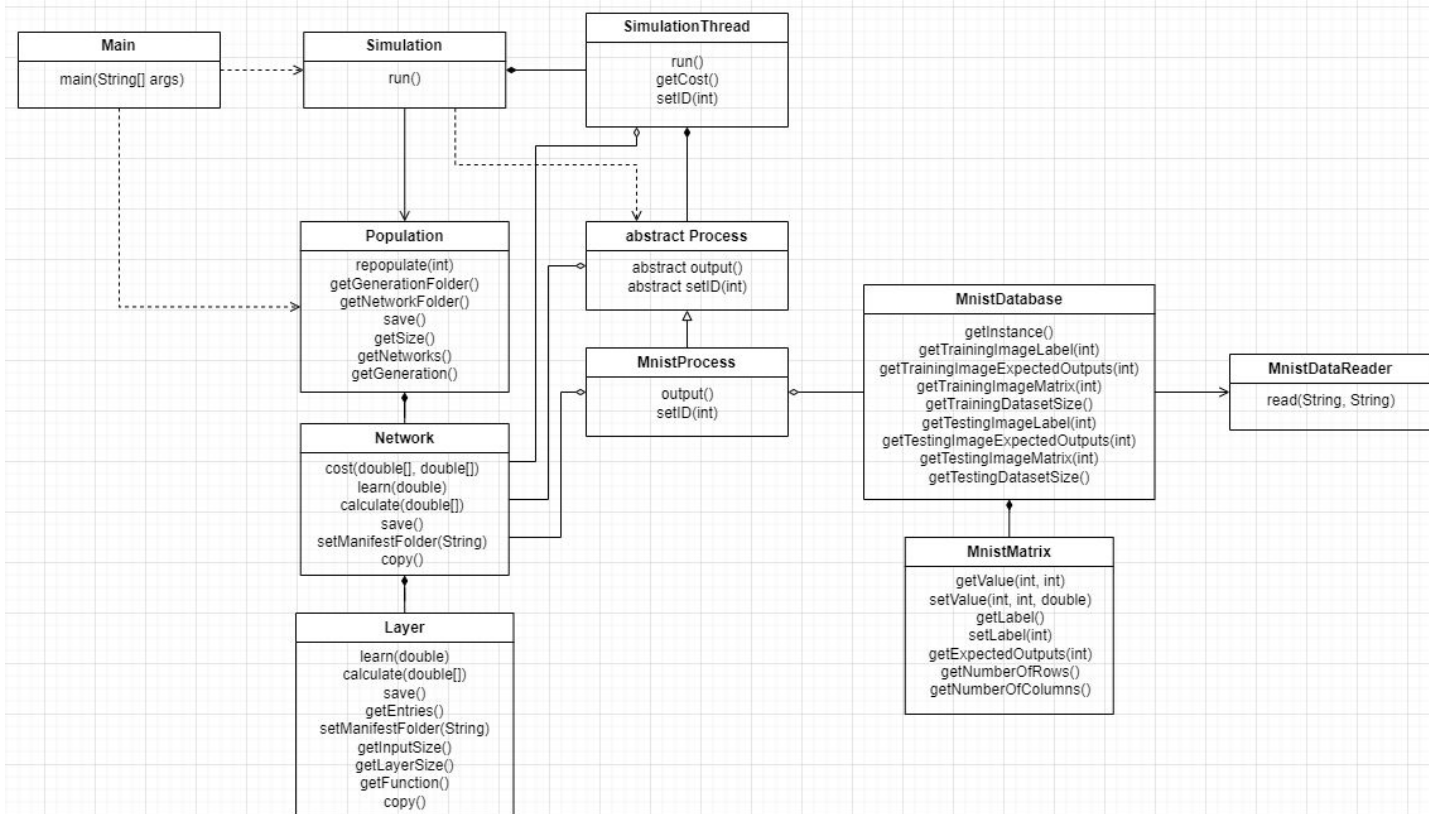
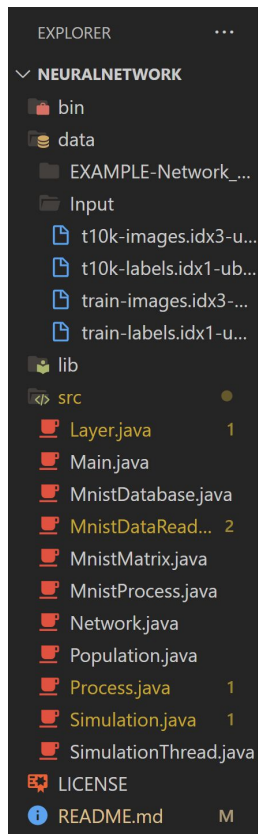


POPULATION BASED TRAINING



CODE STRUCTURE

1859 Lines of Code




▶ **DEMO**



▶ RECAP

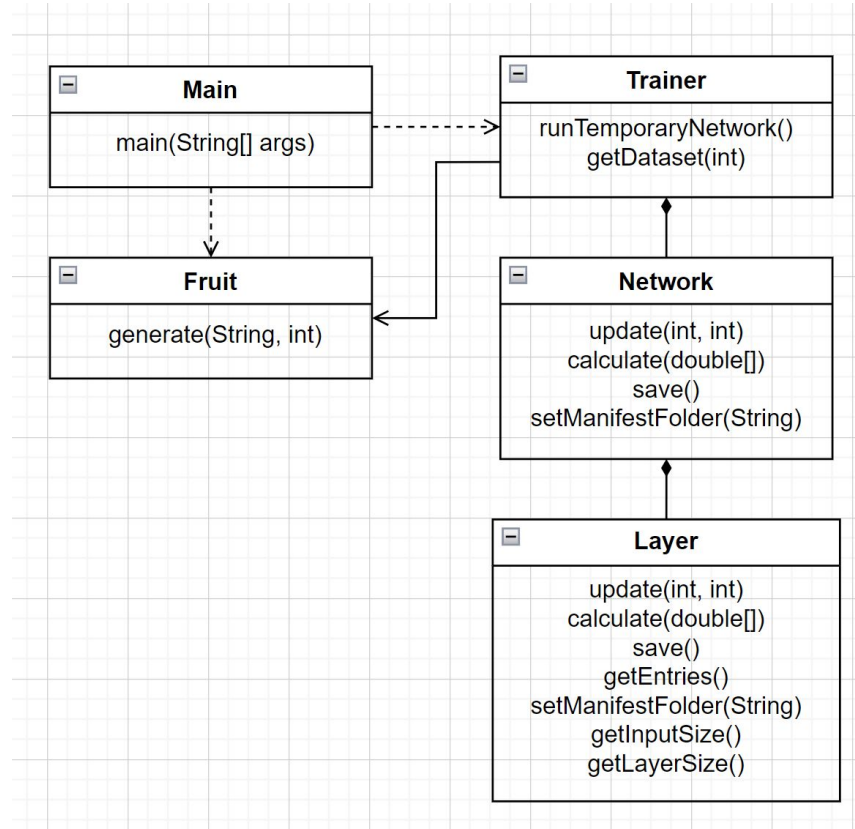
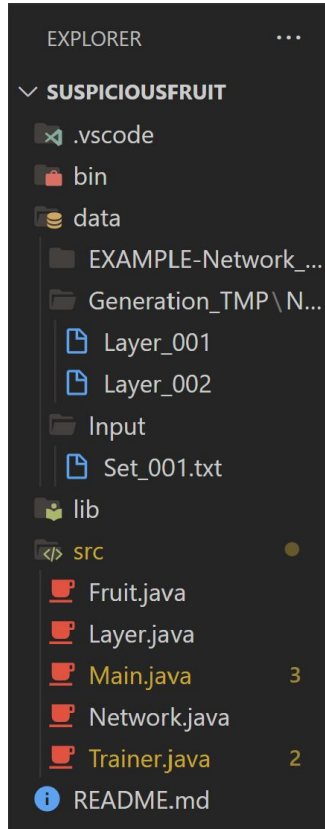
NEURAL NETWORK

- A neural network is composed of many layers with many nodes.
- A node is more "activated" as its value approaches 1.
- Population-Based Training (PBT) provides an efficient way to train a successful neural network model.




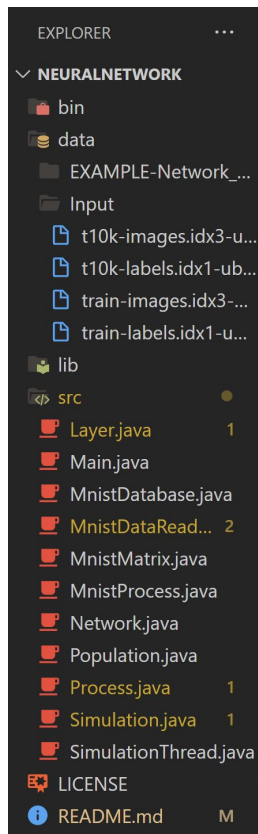
CODE STRUCTURE: FROM THIS

892 Lines of Code



CODE STRUCTURE: TO THIS

1859 Lines of Code




▶ **GOING FORWARD**




THE AI IS SLOW

- Performing completely random changes to the AI is not very efficient...
 - I let my AI run overnight once and it made no progress at all.
- Solution: backpropagation and gradient descent.
 - Calculus concepts that will take at least a good couple weeks to get a grasp of with then another week to flesh out the new structure (not including the actual implementation of it in code).



Computed by WolframAlpha



Computed by WolframAlpha

BEYOND FRUIT AND MNIST

- Face/object recognition
- Checkers





▶ THANK YOU ◀

POST-PRESENTATION

- All code with extensive documentation, links to helpful resources, and a pseudo-dev journal via GitHub commits are found at this repository:
 - <https://github.com/savojosh/NeuralNetwork/tree/MNIST>
 - Code from when I was seeing if a fruit was poisonous or not is not posted at this repository.
- Special thanks to:
 - 3Blue1Brown for the video used in this presentation;
 - Google DeepMind for Population-Based Training (PBT);
 - Sebastian Lague for an intro on how to actually code a neural network;
 - Dr. Cao for help with Java Reflections;
 - Jason Brownlee for extremely informational blog posts;
 - And to all other sources that helped me along the way.